

Sri Sivasubramaniya Nadar College of Engineering, Chennai

(An autonomous Institution affiliated to Anna University)

Degree & Branch	M.Tech (Integrated) Computer Science & Engineering
Semester	V
Subject Code & Name	ICS1512 – Machine Learning Algorithms Laboratory
Academic Year	2025–2026 (Odd)
Batch	2023–2028

Name: I.S.Rajesh

Register No.: 3122237001042

Experiment 4: Ensemble Prediction and Decision Tree Model Evaluation

Aim

To build and evaluate classifiers such as Decision Tree, AdaBoost, Gradient Boosting, XGBoost, Random Forest, and Stacked Models using 5-Fold Cross-Validation and hyperparameter tuning.

Libraries Used

numpy, pandas, sklearn, matplotlib, seaborn, xgboost, ucimlrepo

Objective

Perform ensemble prediction and compare Decision Tree, AdaBoost, Gradient Boosting, XGBoost, Random Forest, and Stacked models, while tuning hyperparameters and analyzing results with cross-validation on Wisconsin Diagnostic Dataset.

Python Code

```
# (i) Import necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, KFold,
    cross_val_score, GridSearchCV
from sklearn.preprocessing import MinMaxScaler, StandardScaler,
    LabelEncoder, OneHotEncoder, Binarizer
```

```

from sklearn.metrics import mean_squared_error,
    root_mean_squared_error, mean_absolute_error, r2_score,
    accuracy_score, precision_score, recall_score, f1_score,
    classification_report, confusion_matrix,
    ConfusionMatrixDisplay, roc_auc_score, roc_curve
from sklearn.svm import SVC
from sklearn.ensemble import AdaBoostClassifier,
    GradientBoostingClassifier, RandomForestClassifier,
    StackingClassifier
from sklearn.tree import DecisionTreeClassifier
from xgboost import XGBClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier

from ucimlrepo import fetch_ucirepo

# (ii) Load dataset
breast_cancer_wisconsin_diagnostic = fetch_ucirepo(id=17)

X = breast_cancer_wisconsin_diagnostic.data.features
y = breast_cancer_wisconsin_diagnostic.data.targets

df = pd.concat([X, y], axis=1)

df.to_csv("breast_cancer_wisconsin_diagnostic.csv", index=False)

target = 'Diagnosis'

# (iii) EDA and Preprocessing

def is_normal(series):
    skew = series.skew()
    return -0.5 <= skew <= 0.5

def has_outliers(series):
    Q1 = series.quantile(0.25)
    Q3 = series.quantile(0.75)
    IQR = Q3 - Q1
    outliers = ((series < (Q1 - 1.5 * IQR)) | (series > (Q3 + 1.5
        * IQR))).sum()
    return outliers > 0

# Separate types
numerical_cols = df.select_dtypes(include=['int64', 'float64']).
    columns.tolist()
categorical_cols = df.select_dtypes(include=['object', 'category'
    ]).columns.tolist()
numerical_cols = [col for col in numerical_cols if col != target]

```

```

# --- Missing Values Handling ---
for col in df.columns:
    if df[col].isnull().sum() > 0:
        if col in numerical_cols:
            if is_normal(df[col]):
                if has_outliers(df[col]):
                    df[col].fillna(df[col].median(), inplace=True)
                else:
                    df[col].fillna(df[col].mean(), inplace=True)
            else:
                df[col].fillna(df[col].median(), inplace=True)
        else:
            df[col].fillna(df[col].mode()[0], inplace=True) # categorical

# --- Outlier Replacing (for numerical columns only) ---
for col in numerical_cols:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    if is_normal(df[col]):
        replacement_value = df[col].mean()
    else:
        replacement_value = df[col].median()

    # Replace outliers
    df[col] = np.where((df[col] < lower_bound) | (df[col] >
        upper_bound), replacement_value, df[col])

# --- Drop rows where target is missing ---
df.dropna(subset=[target], inplace=True)

# --- Feature Engineering (if needed) ---
df["circularity1"] = 4 * np.pi * df["area1"] / (df["perimeter1"]
    ** 2 + 1e-6)
df["area_per_radius1"] = df["area1"] / (df["radius1"] + 1e-6)

# --- Encoding categorical features ---
is_classification = True

if is_classification:
    le = LabelEncoder()
    df[target] = le.fit_transform(df[target]) # encode target
    for col in categorical_cols:
        df[col] = le.fit_transform(df[col]) # label encoding for
        classification
else:

```

```

# Regression: target guided ordinal encoding
for col in categorical_cols:
    ordered_labels = df.groupby(col)[target].mean().
        sort_values().index
    mapping = {k: i for i, k in enumerate(ordered_labels)}
    df[col] = df[col].map(mapping)

# --- Histogram Subplots ---
numerical_cols = df.select_dtypes(include=['int64', 'float64']).
    columns.tolist()
categorical_cols = df.select_dtypes(include=['object', 'category'
]).columns.tolist()
numerical_cols = [col for col in numerical_cols if col != target]
n_cols = 5 # Number of plots per row
n_rows = int(np.ceil(len(numerical_cols) / n_cols))

fig, axes = plt.subplots(n_rows, n_cols, figsize=(20, 4 * n_rows)
    )
axes = axes.flatten()

for i, col in enumerate(numerical_cols):
    df[col].hist(ax=axes[i], bins=30)
    axes[i].set_title(col)

# Turn off unused subplots
for j in range(i+1, len(axes)):
    axes[j].axis('off')

fig.suptitle("Histogram of Features", fontsize=20)
plt.tight_layout()
plt.show()

# --- Boxplot Subplots ---
fig, axes = plt.subplots(n_rows, n_cols, figsize=(20, 4 * n_rows)
    )
axes = axes.flatten()

for i, col in enumerate(numerical_cols):
    axes[i].boxplot(df[col])
    axes[i].set_title(col)

for j in range(i+1, len(axes)):
    axes[j].axis('off')

fig.suptitle("Boxplot for Outlier Detection", fontsize=20)
plt.tight_layout()
plt.show()

# --- Correlation Heatmap ---
plt.figure(figsize=(20, 16))

```

```

sns.heatmap(df.corr(), annot=True, fmt=".2f", cmap='coolwarm',
            linewidths=0.5)
plt.title("Feature Correlation Heatmap", fontsize=18)
plt.xticks(rotation=45, ha='right', fontsize=8)
plt.yticks(fontsize=8)
plt.tight_layout()
plt.show()

# (iv) Splitting dataset

X = df.drop(columns=[target])
y = df[target]

# Splitting: Train (60%), Validation (20%), Test (20%)
X_train, X_temp, y_train, y_temp = train_test_split(X, y,
            test_size=0.4, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp,
            test_size=0.5, random_state=42)

print(f'Train: {X_train.shape}, Validation: {X_val.shape}, Test:
      {X_test.shape}')

# Store results
best_models = {}
y_val_preds = {}
y_test_preds = {}

# Common scoring metric
scoring = 'accuracy'

# ----- Decision Tree -----
param_dt = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 5, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4, 5, 10]
}
grid_dt = GridSearchCV(
    DecisionTreeClassifier(random_state=42),
    param_dt, cv=5, scoring=scoring, n_jobs=-1
)
grid_dt.fit(X_train, y_train)
model_dt = grid_dt.best_estimator_
y_val_dt = model_dt.predict(X_val)
y_test_dt = model_dt.predict(X_test)

best_models['DecisionTree'] = model_dt
y_val_preds['DecisionTree'] = y_val_dt
y_test_preds['DecisionTree'] = y_test_dt

# ----- Random Forest -----

```

```

param_rf = {
    'n_estimators': [100, 200],
    'max_depth': [None, 5, 10, 20],
    'criterion': ['gini', 'entropy'],
    'min_samples_split': [2, 5, 10],
    'max_features': ['sqrt', 'log2', None, 0.3, 0.5, 0.7]
}
grid_rf = GridSearchCV(
    RandomForestClassifier(random_state=42),
    param_rf, cv=5, scoring=scoring, n_jobs=-1
)
grid_rf.fit(X_train, y_train)
model_rf = grid_rf.best_estimator_
y_val_rf = model_rf.predict(X_val)
y_test_rf = model_rf.predict(X_test)

best_models['RandomForest'] = model_rf
y_val_preds['RandomForest'] = y_val_rf
y_test_preds['RandomForest'] = y_test_rf

# ----- AdaBoost -----
param_ab = {
    'n_estimators': [50, 100, 200],
    'learning_rate': [0.01, 0.1, 1],
    'estimator': [
        DecisionTreeClassifier(max_depth=1),
        DecisionTreeClassifier(max_depth=2),
        DecisionTreeClassifier(max_depth=3)
    ]
}
grid_ab = GridSearchCV(
    AdaBoostClassifier(random_state=42),
    param_ab, cv=5, scoring=scoring, n_jobs=-1
)
grid_ab.fit(X_train, y_train)
model_ab = grid_ab.best_estimator_
y_val_ab = model_ab.predict(X_val)
y_test_ab = model_ab.predict(X_test)

best_models['AdaBoost'] = model_ab
y_val_preds['AdaBoost'] = y_val_ab
y_test_preds['AdaBoost'] = y_test_ab

# ----- Gradient Boost -----
param_gb = {
    'n_estimators': [100, 200],
    'learning_rate': [0.01, 0.1, 0.2],
    'max_depth': [3, 5, 7],
    'subsample': [0.6, 0.8, 1.0]
}

```

```

grid_gb = GridSearchCV(
    GradientBoostingClassifier(random_state=42),
    param_gb, cv=5, scoring=scoring, n_jobs=-1
)
grid_gb.fit(X_train, y_train)
model_gb = grid_gb.best_estimator_
y_val_gb = model_gb.predict(X_val)
y_test_gb = model_gb.predict(X_test)

best_models['GradientBoost'] = model_gb
y_val_preds['GradientBoost'] = y_val_gb
y_test_preds['GradientBoost'] = y_test_gb

# ----- XGBoost -----
param_xgb = {
    'n_estimators': [100, 200],
    'learning_rate': [0.01, 0.1, 0.2],
    'max_depth': [3, 5, 7],
    'subsample': [0.8, 1.0],
    'colsample_bytree': [0.8, 1.0],
    'gamma': [0, 0.1, 0.5, 1, 5]
}
grid_xgb = GridSearchCV(
    XGBClassifier(use_label_encoder=False, eval_metric='logloss',
        random_state=42),
    param_xgb, cv=5, scoring=scoring, n_jobs=-1
)
grid_xgb.fit(X_train, y_train)
model_xgb = grid_xgb.best_estimator_
y_val_xgb = model_xgb.predict(X_val)
y_test_xgb = model_xgb.predict(X_test)

best_models['XGBoost'] = model_xgb
y_val_preds['XGBoost'] = y_val_xgb
y_test_preds['XGBoost'] = y_test_xgb

def tune_svm(X_train, y_train):
    pipe = Pipeline([
        ("scaler", StandardScaler()),
        ("svm", SVC(probability=True, random_state=42))
    ])

    param_grid = {
        "svm__kernel": ["linear", "rbf", "poly"],
        "svm__C": [0.1, 1, 10],
        "svm__gamma": ["scale", "auto"]
    }

    grid = GridSearchCV(pipe, param_grid, cv=5, scoring="accuracy", n_jobs=-1)
    grid.fit(X_train, y_train)

```

```

    print("Best SVM Params:", grid.best_params_)
    return grid.best_estimator_

best_svm = tune_svm(X_train, y_train)
nb_model = GaussianNB()
dt_model = best_models['DecisionTree'] # reuse tuned one
rf_model = best_models['RandomForest'] # reuse tuned one
knn_model = Pipeline([("scaler", StandardScaler()), ("knn",
    KNeighborsClassifier())])

stack1 = StackingClassifier(
    estimators=[("svm", best_svm), ("nb", nb_model), ("dt",
        dt_model)],
    final_estimator=LogisticRegression(max_iter=500, random_state
        =42),
    cv=5,
    n_jobs=-1
)

stack2 = StackingClassifier(
    estimators=[("svm", best_svm), ("nb", nb_model), ("dt",
        dt_model)],
    final_estimator=RandomForestClassifier(n_estimators=200,
        random_state=42),
    cv=5,
    n_jobs=-1
)

stack3 = StackingClassifier(
    estimators=[("svm", best_svm), ("dt", dt_model), ("knn",
        knn_model)],
    final_estimator=LogisticRegression(max_iter=500, random_state
        =42),
    cv=5,
    n_jobs=-1
)

for name, model in {
    "Stacked_LogReg": stack1,
    "Stacked_RF": stack2,
    "Stacked_LogReg_KNN": stack3
}.items():
    model.fit(X_train, y_train)
    y_val_preds[name] = model.predict(X_val)
    y_test_preds[name] = model.predict(X_test)
    best_models[name] = model

for i in best_models:
    print(i, best_models[i].get_params(), '\n')

# (vi) Evaluation

```



```

# Evaluating Model using Performance Metrics

def evaluate_model(y_true, y_pred, is_classification, X, model,
dataset_name):
    print(f"\n Evaluation - {dataset_name}")
    if is_classification:
        print("Accuracy :", round(accuracy_score(y_true, y_pred),
4))
        print("Precision:", round(precision_score(y_true, y_pred,
average='weighted'), 4))
        print("Recall    :", round(recall_score(y_true, y_pred,
average='weighted'), 4))
        print("F1 Score :", round(f1_score(y_true, y_pred,
average='weighted'), 4))
        print("\nClassification Report:\n", classification_report
(y_true, y_pred))

    # ROC Curve: Only for binary classification
    if len(np.unique(y_true)) == 2 and model is not None and
hasattr(model, "predict_proba"):
        y_probs = model.predict_proba(X)[: , 1]
        fpr, tpr, _ = roc_curve(y_true, y_probs)
        auc_score = roc_auc_score(y_true, y_probs)
        print("ROC AUC Score:", round(auc_score, 4))

        # Plot ROC
        plt.figure(figsize=(6, 4))
        plt.plot(fpr, tpr, label=f"AUC = {auc_score:.4f}")
        plt.plot([0, 1], [0, 1], 'k--', label='Random Guess')
        plt.xlabel("False Positive Rate")
        plt.ylabel("True Positive Rate")
        plt.title(f"ROC Curve - {dataset_name}")
        plt.legend()
        plt.grid(True)
        plt.tight_layout()
        plt.show()
    else:
        n, p = X.shape
        r2 = r2_score(y_true, y_pred)
        adjusted_r2 = 1 - (1 - r2) * (n - 1) / (n - p - 1)
        print("Mean Squared Error:", mean_squared_error(y_true,
y_pred))
        print("Root Mean Squared Error:", root_mean_squared_error
(y_true, y_pred))
        print("Mean Absolute Error:", mean_absolute_error(y_true,
y_pred))
        print("R2 Score:", r2)
        print("Adjusted R2 Score:", adjusted_r2)

for model_name, model in best_models.items():

```

```

print(f"\n{'='*20} {model_name} - Validation {'='*20}")
evaluate_model(y_val, y_val_preds[model_name], True, X_val,
               model, f"{model_name} (Validation)")

print(f"\n{'='*20} {model_name} - Test {'='*20}")
evaluate_model(y_test, y_test_preds[model_name], True, X_test,
               model, f"{model_name} (Test)")

# Evaluating Model on Test and Validation Sets (Without
  Performance Metrics)

def plot_actual_vs_predicted(y_true, y_pred, title):
    plt.figure(figsize=(6, 4))
    plt.scatter(y_true, y_pred, alpha=0.5, edgecolor='k')
    plt.plot([y_true.min(), y_true.max()], [y_true.min(), y_true.
        max()], 'r--')
    plt.xlabel("Actual")
    plt.ylabel("Predicted")
    plt.title(f"Actual vs Predicted - {title}")
    plt.grid(True)
    plt.tight_layout()
    plt.show()

def plot_residuals(y_true, y_pred, title):
    residuals = y_true - y_pred
    plt.figure(figsize=(6, 4))
    plt.scatter(y_pred, residuals, alpha=0.5, edgecolor='k')
    plt.axhline(y=0, color='red', linestyle='--')
    plt.xlabel("Predicted")
    plt.ylabel("Residuals")
    plt.title(f"Residual Plot - {title}")
    plt.grid(True)
    plt.tight_layout()
    plt.show()

def plot_residual_distribution(y_true, y_pred, title):
    residuals = y_true - y_pred
    plt.figure(figsize=(6, 4))
    sns.histplot(residuals, kde=True, color='skyblue')
    plt.title(f"Residual Distribution - {title}")
    plt.xlabel("Residuals")
    plt.ylabel("Frequency")
    plt.grid(True)
    plt.tight_layout()
    plt.show()

def plot_confusion_matrix(y_true, y_pred, title):
    ConfusionMatrixDisplay.from_predictions(y_true, y_pred)
    plt.title(title)
    plt.show()

```

```

for model_name in best_models.keys():
    print(f"\nVisualization for {model_name}")

    y_val_pred = y_val_preds[model_name]
    y_test_pred = y_test_preds[model_name]

    if not is_classification:
        # Validation plots
        plot_actual_vs_predicted(y_val, y_val_pred, f"{model_name} - Validation Set")
        plot_residuals(y_val, y_val_pred, f"{model_name} - Validation Set")
        plot_residual_distribution(y_val, y_val_pred, f"{model_name} - Validation Set")

        # Test plots
        plot_actual_vs_predicted(y_test, y_test_pred, f"{model_name} - Test Set")
        plot_residuals(y_test, y_test_pred, f"{model_name} - Test Set")
        plot_residual_distribution(y_test, y_test_pred, f"{model_name} - Test Set")

    else:
        # Confusion matrices
        plot_confusion_matrix(y_val, y_val_pred, f"{model_name} - Validation Set")
        plot_confusion_matrix(y_test, y_test_pred, f"{model_name} - Test Set")

# (vii) K-Fold Cross Validation

kfold = KFold(n_splits=5, shuffle=True, random_state=42)
if is_classification:
    score = 'accuracy'
else:
    score = 'r2'

cv_results = {}

for model_name, model in best_models.items():
    print(f"\nCross-Validation for {model_name}")

    scores = cross_val_score(model, X, y, cv=kfold, scoring=score)

    cv_results[model_name] = scores

    print("Cross Validation Scores:", scores)
    print("Average CV Score:", round(scores.mean(), 4))

```

Output Screenshots

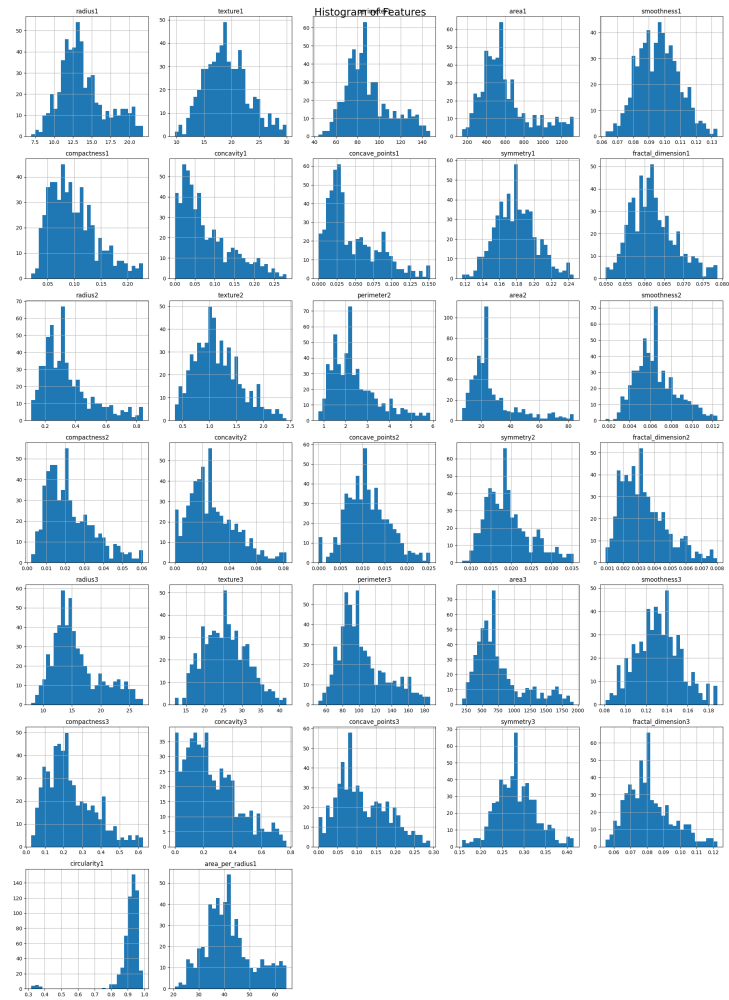


Figure 1: Feature Distribution

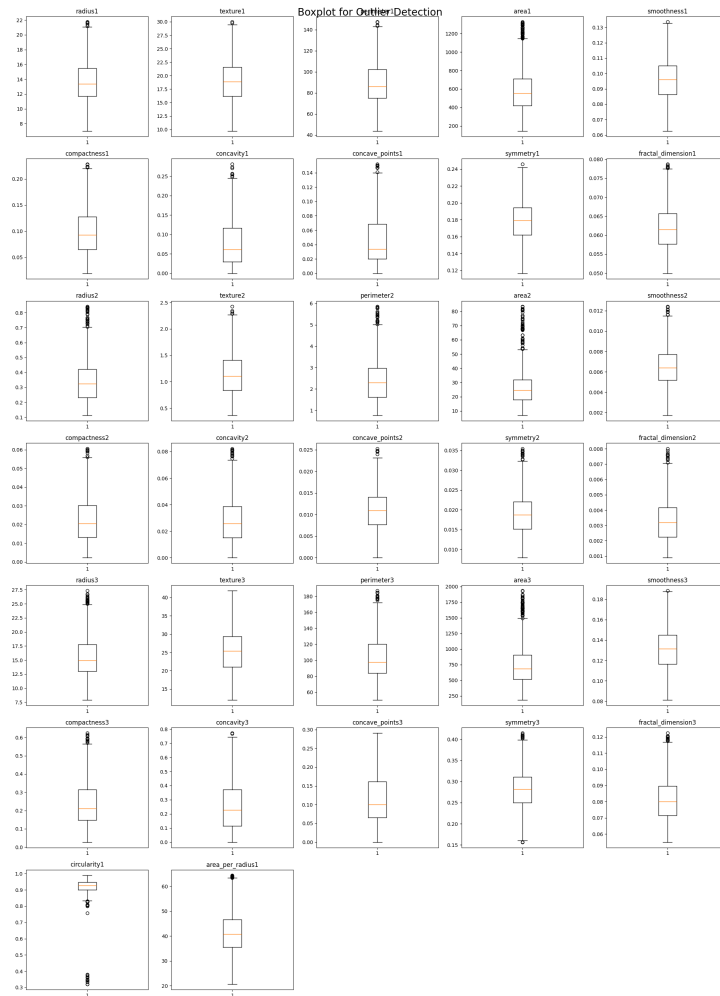


Figure 2: Feature Outliers

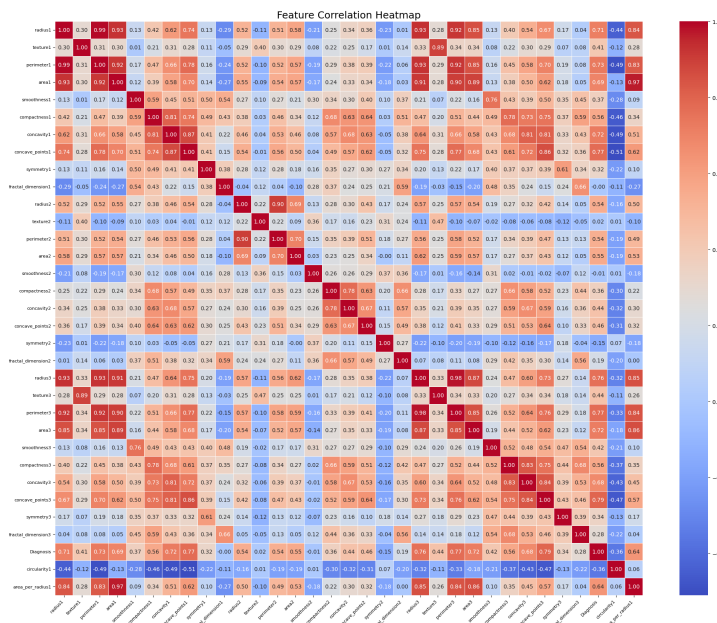


Figure 3: Feature Importance (Correlation Heatmap)

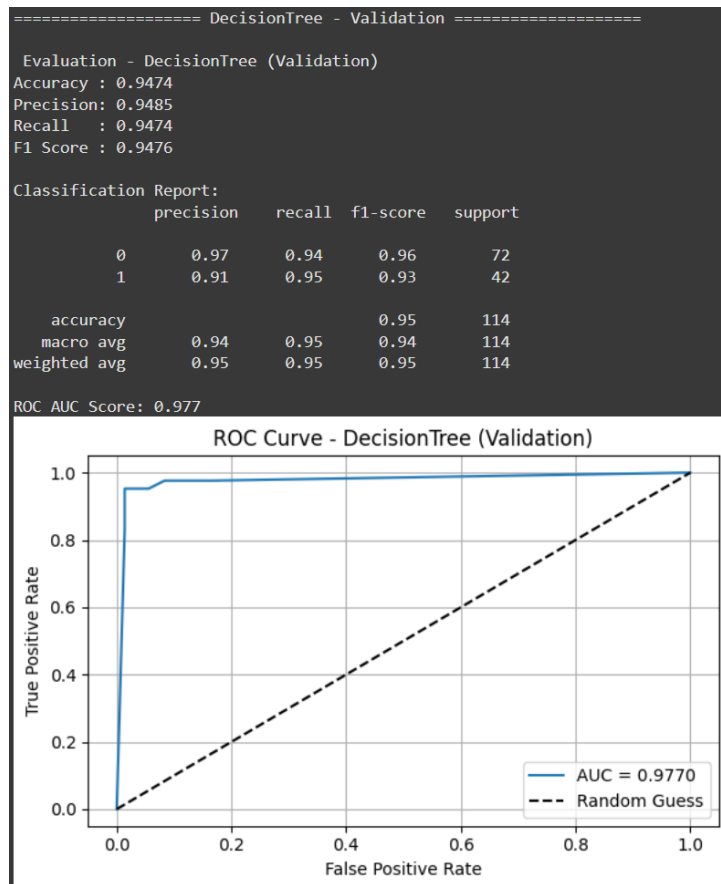


Figure 4: Decision Tree Validation Set Performance

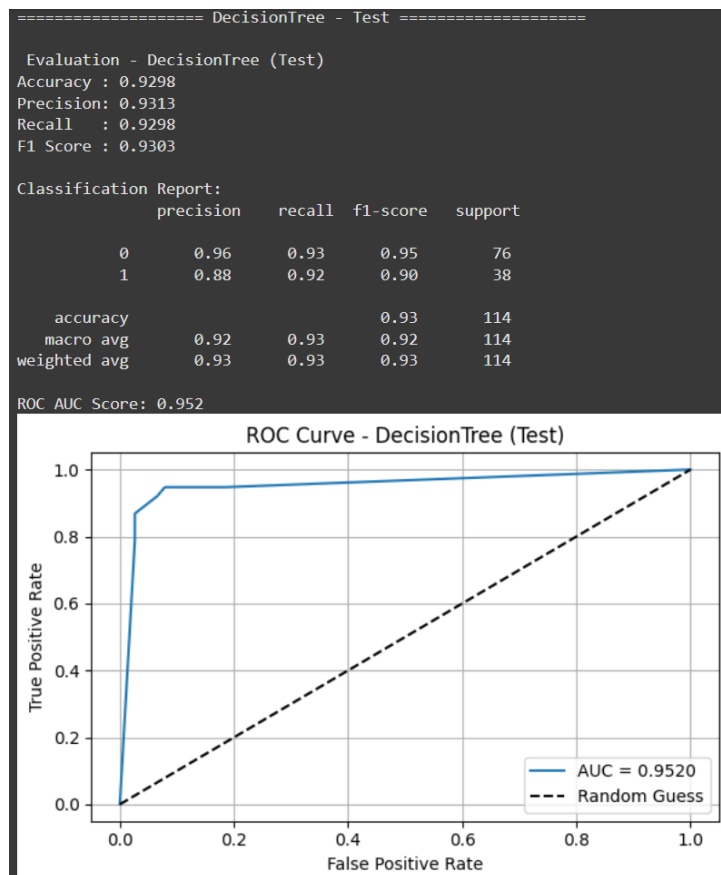


Figure 5: Decision Tree Test Set Performance

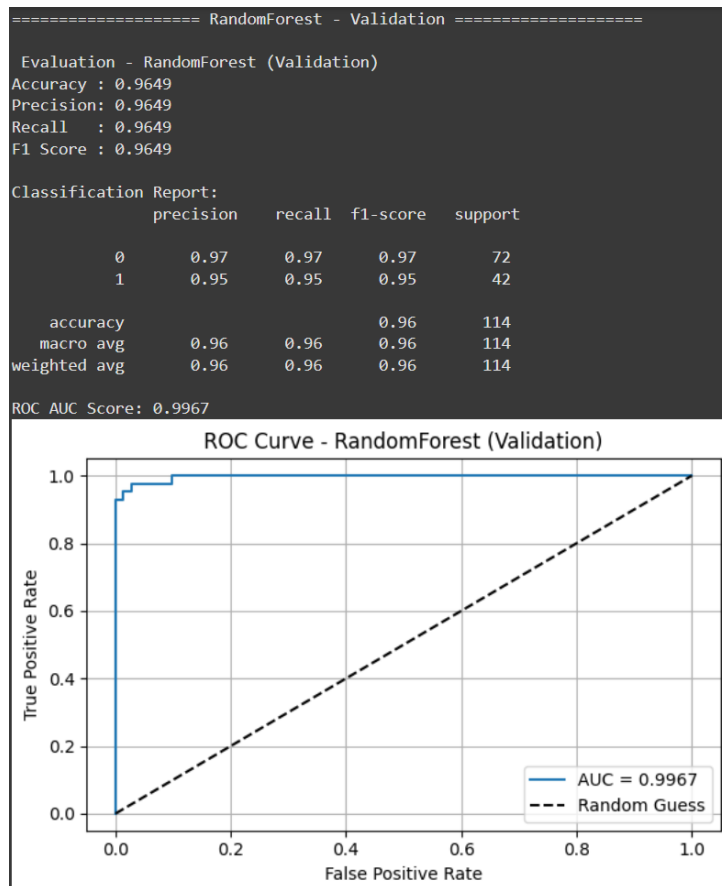


Figure 6: Random Forest Validation Set Performance


```

===== RandomForest - Test =====

Evaluation - RandomForest (Test)
Accuracy : 0.9561
Precision: 0.956
Recall   : 0.9561
F1 Score : 0.956

Classification Report:

```

	precision	recall	f1-score	support
0	0.96	0.97	0.97	76
1	0.95	0.92	0.93	38
accuracy			0.96	114
macro avg	0.95	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

```

ROC AUC Score: 0.9962

```

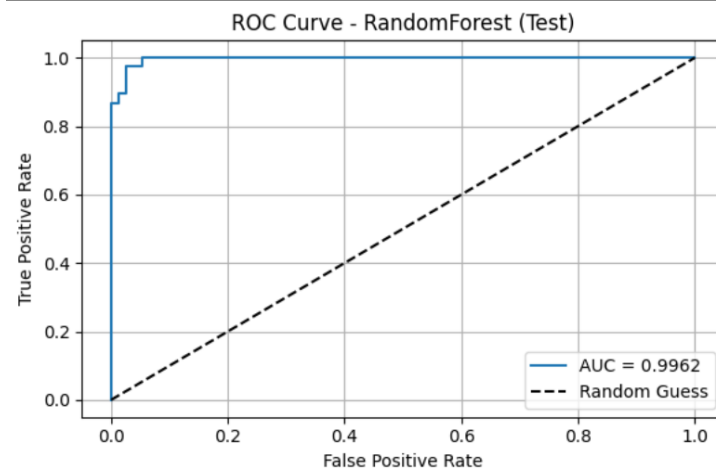


Figure 7: Random Forest Test Set Performance

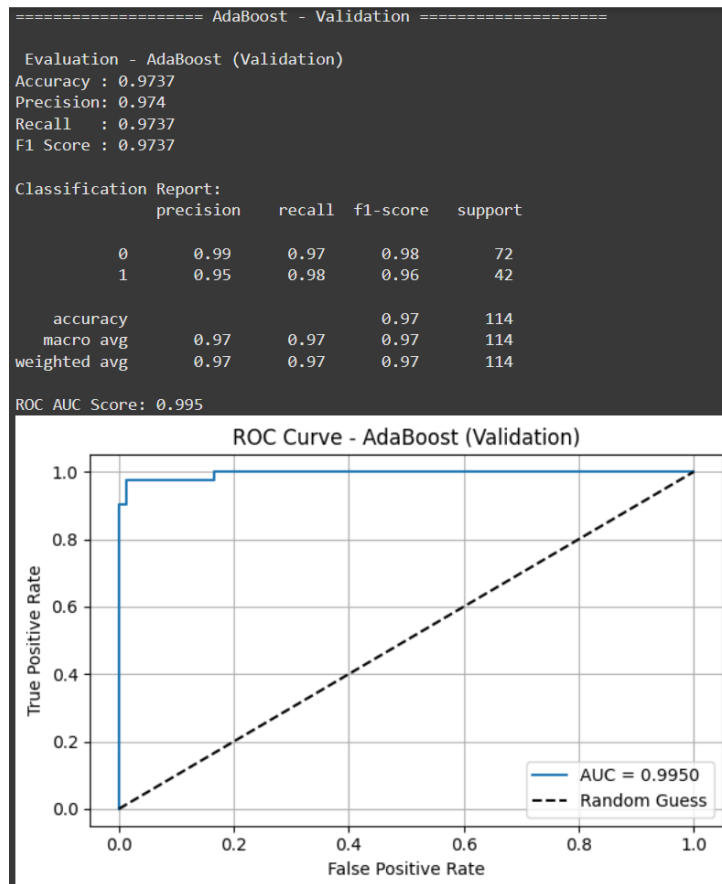


Figure 8: AdaBoost Validation Set Performance

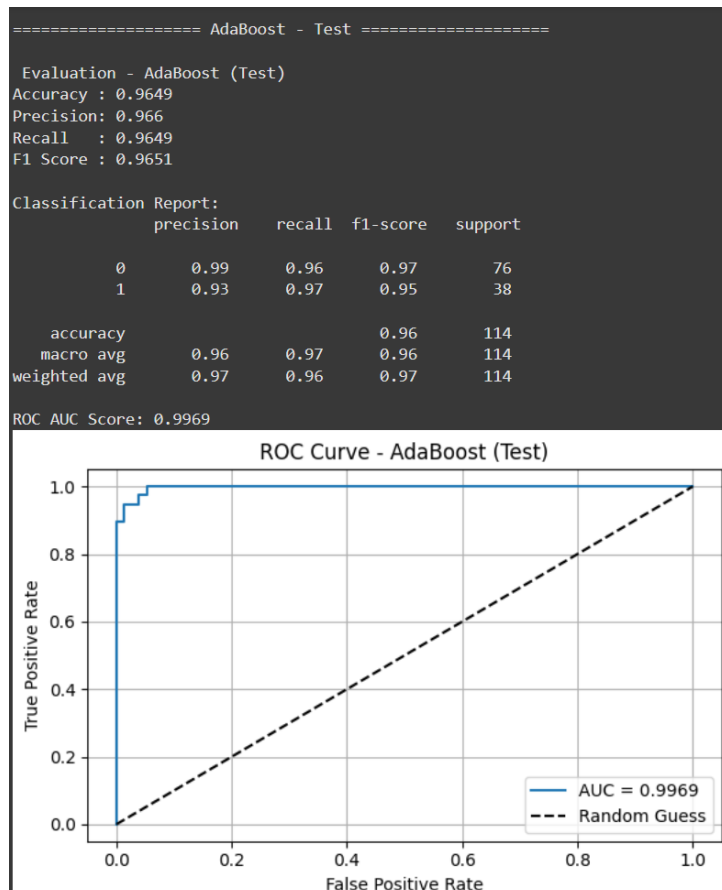


Figure 9: AdaBoost Test Set Performance

```

===== GradientBoost - Validation =====

Evaluation - GradientBoost (Validation)
Accuracy : 0.9737
Precision: 0.9737
Recall   : 0.9737
F1 Score : 0.9736

Classification Report:
      precision    recall  f1-score   support

     0       0.97       0.99       0.98        72
     1       0.98       0.95       0.96        42

   accuracy       0.97               114
  macro avg       0.97       0.97       0.97       114
weighted avg       0.97       0.97       0.97       114

ROC AUC Score: 0.996

```

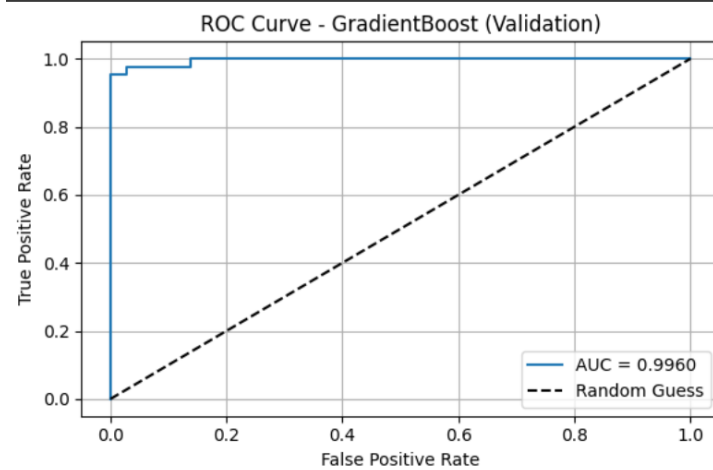


Figure 10: Gradient Boost Validation Set Performance

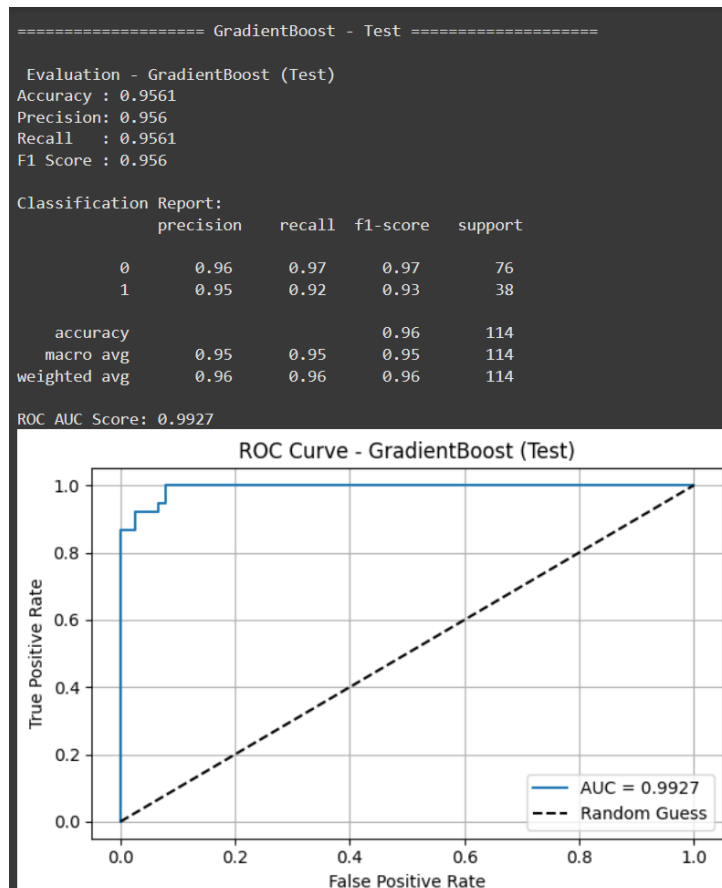


Figure 11: Gradient Boost Test Set Performance

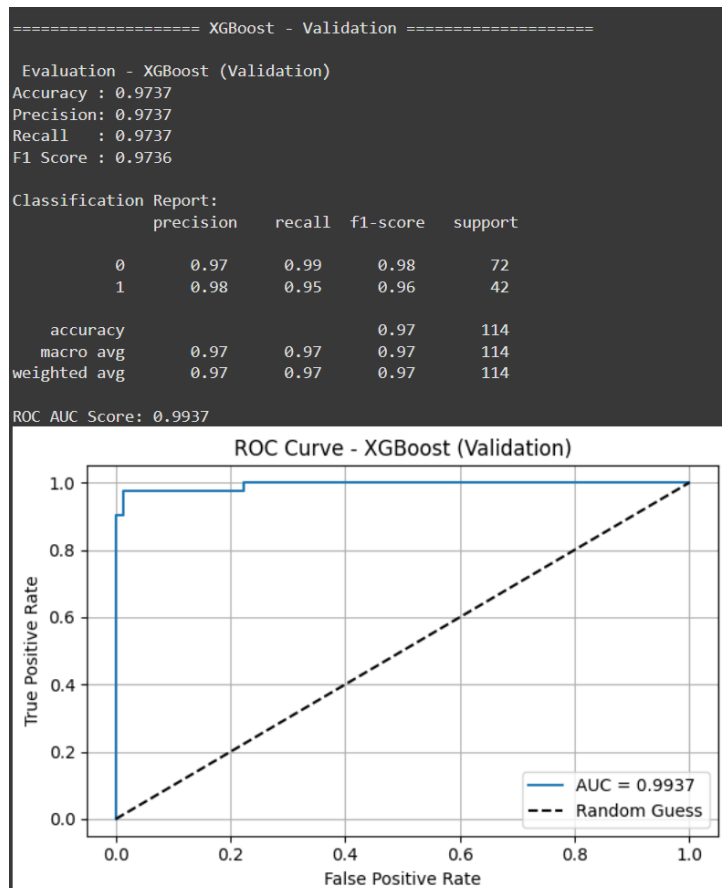


Figure 12: XGBoost Validation Set Performance

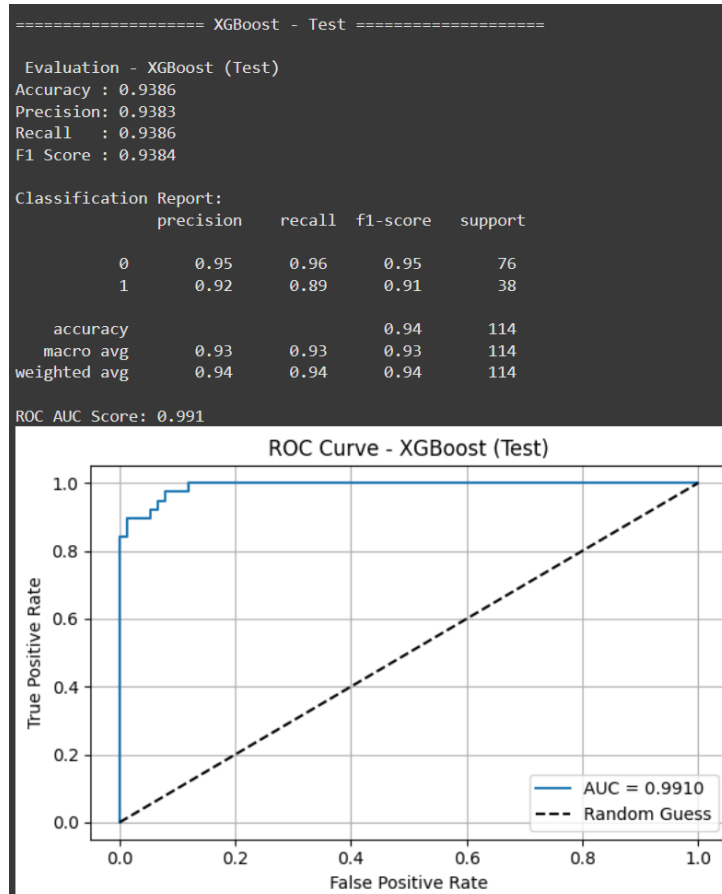


Figure 13: XGBoost Test Set Performance

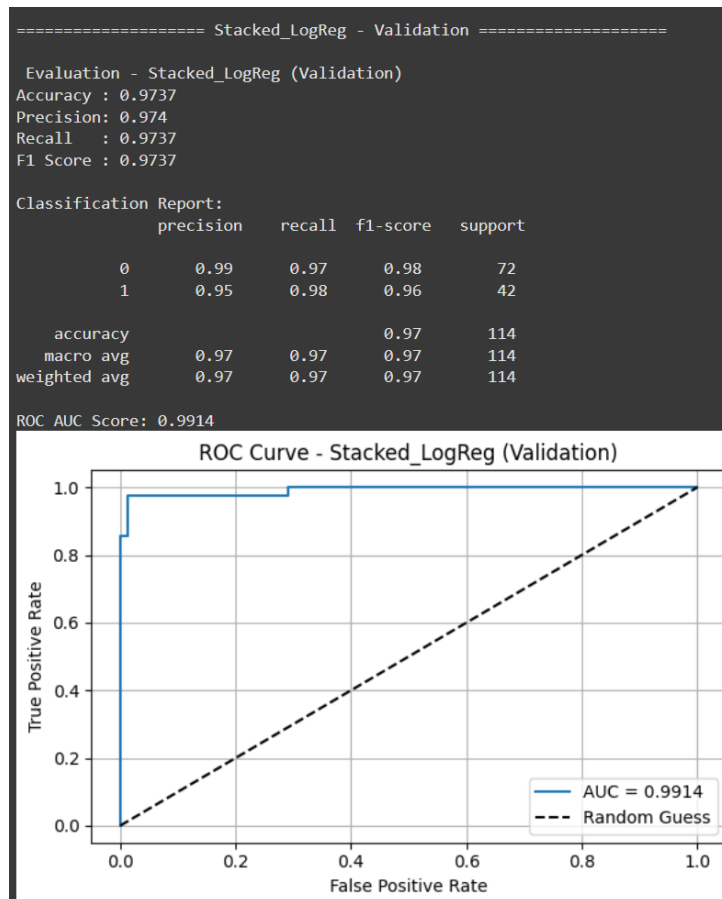


Figure 14: Stacked Ensemble 1 Validation Set Performance

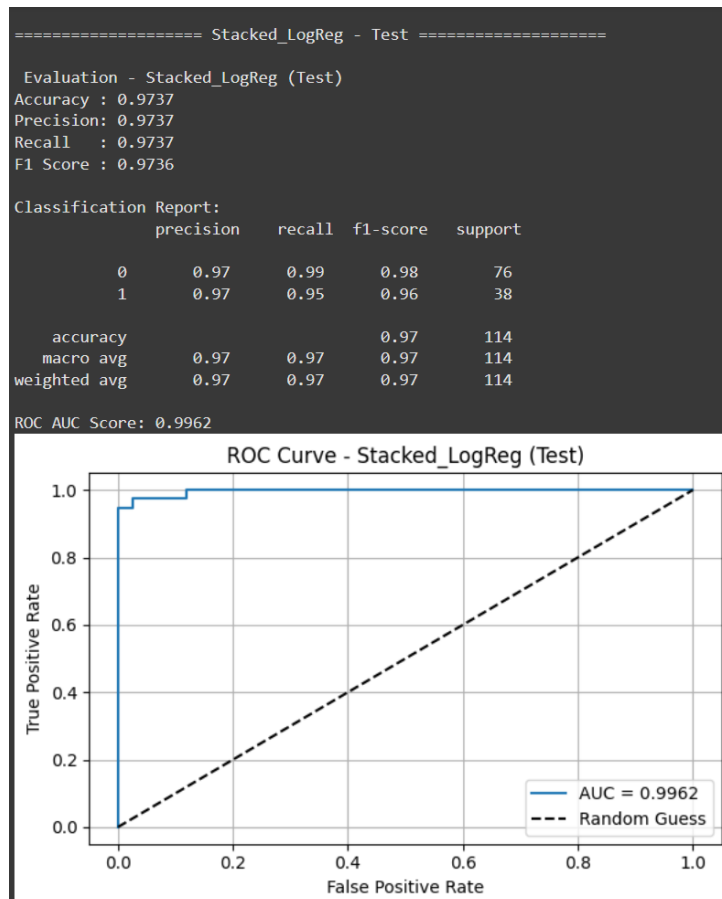


Figure 15: Stacked Ensemble 1 Test Set Performance

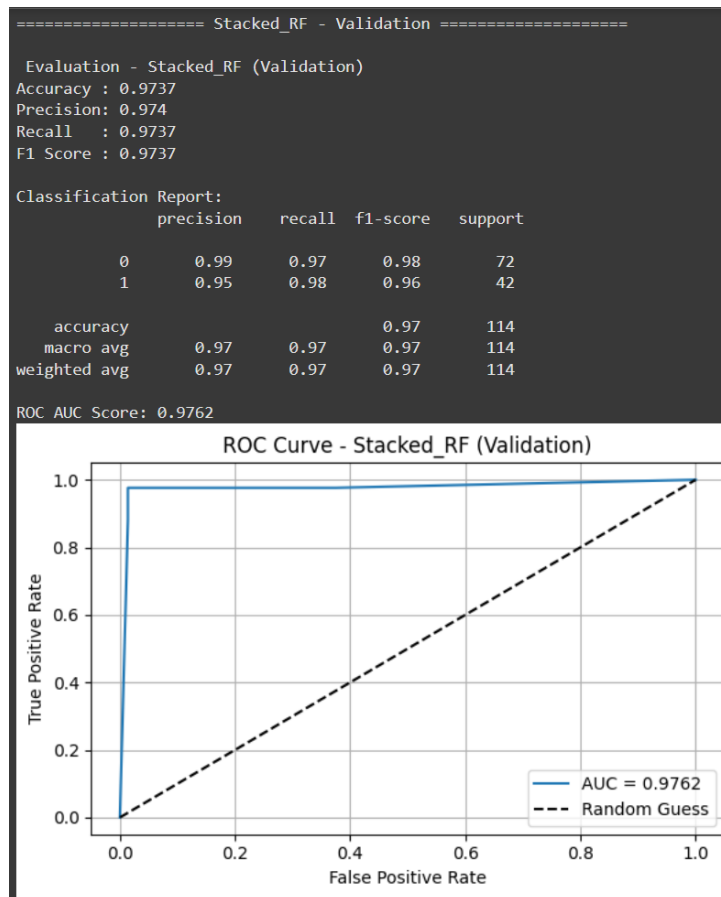


Figure 16: Stacked Ensemble 2 Validation Set Performance

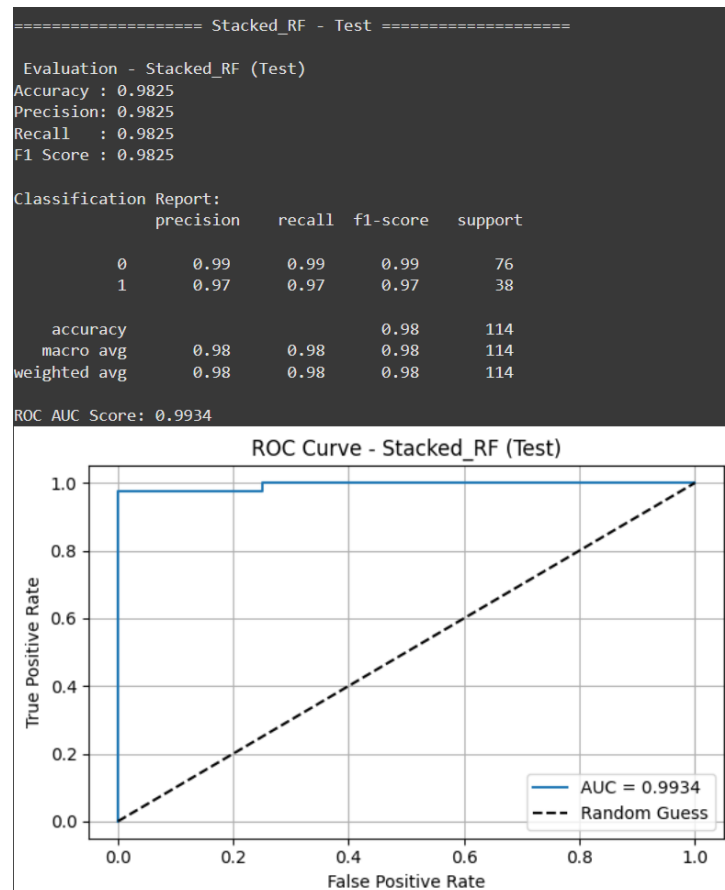


Figure 17: Stacked Ensemble 2 Test Set Performance

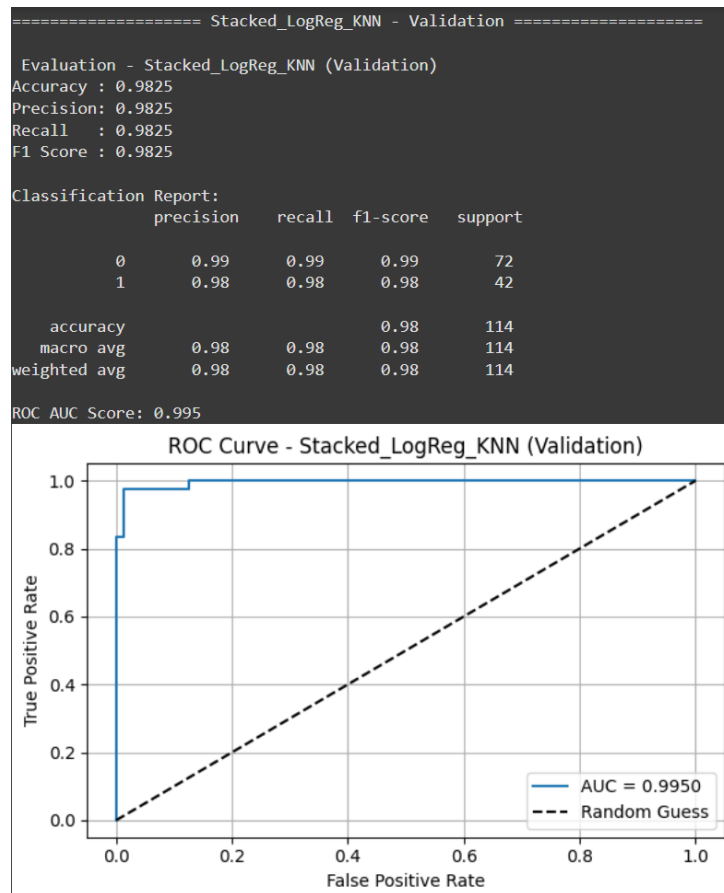


Figure 18: Stacked Ensemble 3 Validation Set Performance

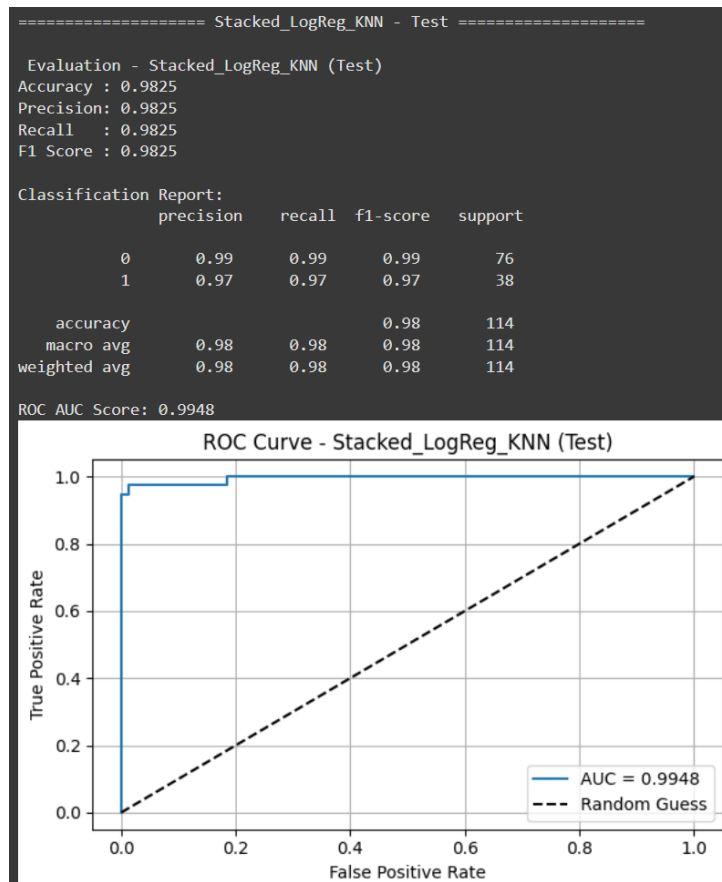


Figure 19: Stacked Ensemble 3 Test Set Performance

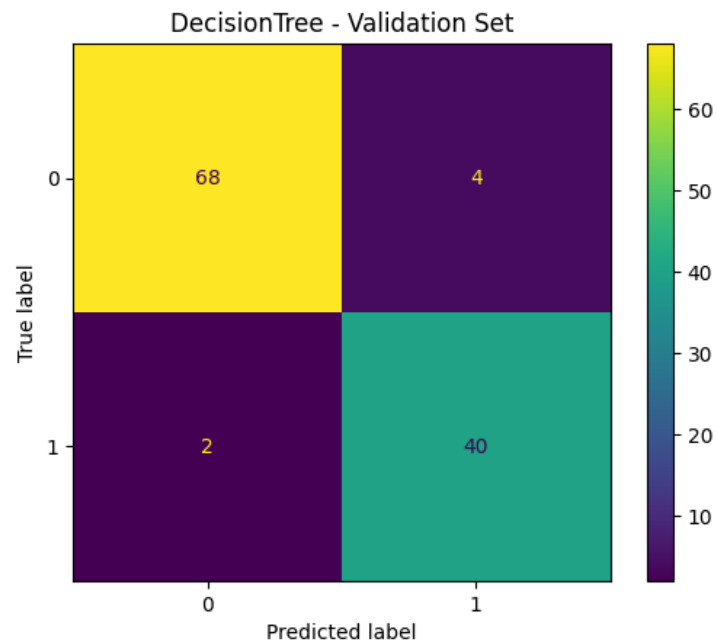


Figure 20: Decision Tree Validation Set Confusion Matrix

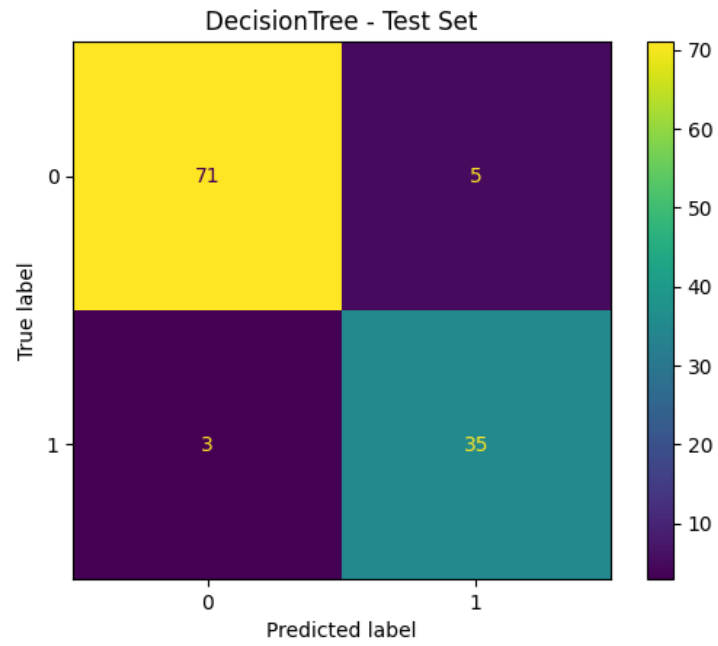


Figure 21: Decision Tree Test Set Confusion Matrix

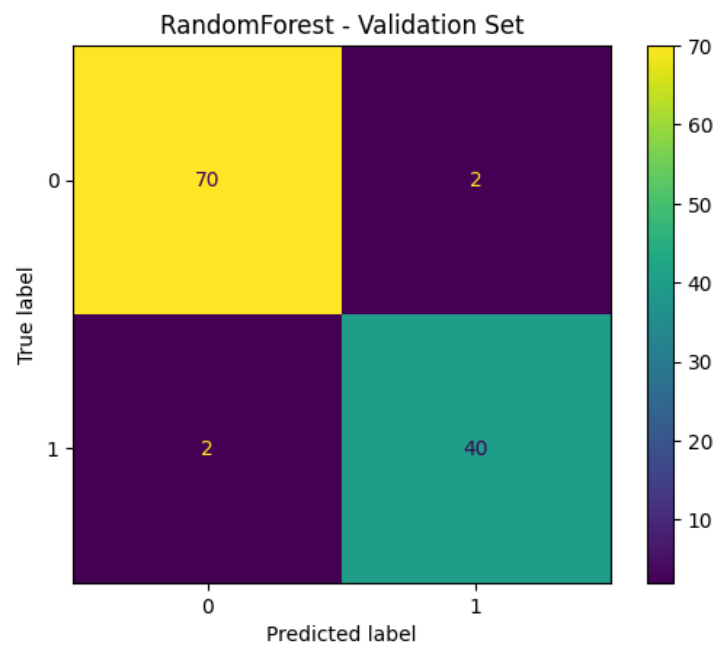


Figure 22: Random Forest Validation Set Confusion Matrix

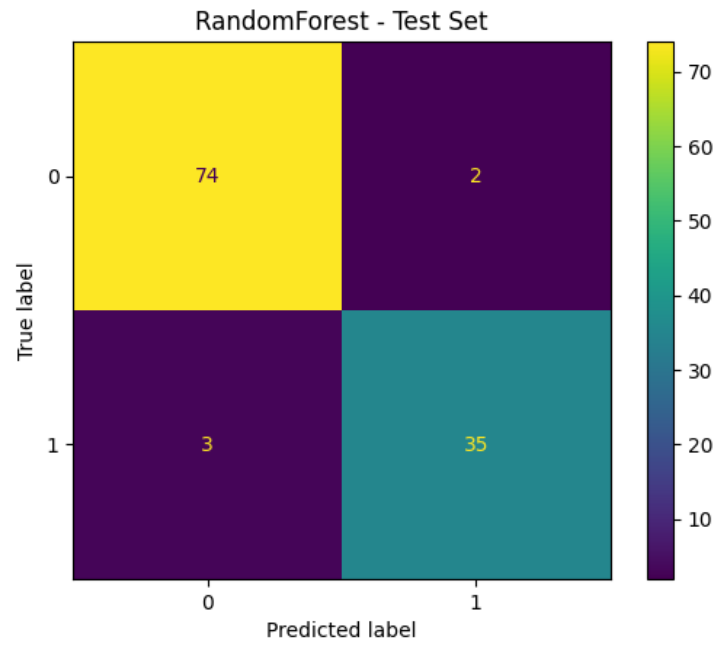


Figure 23: Random Forest Test Set Confusion Matrix

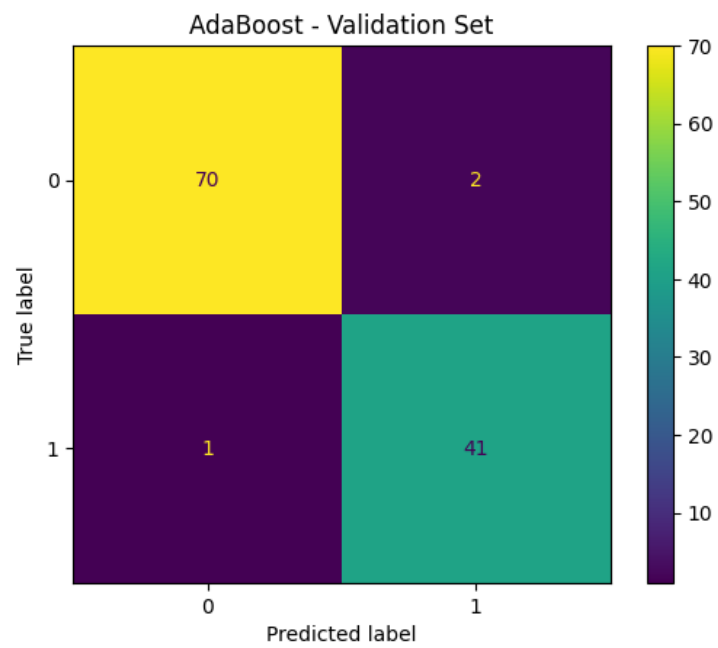


Figure 24: AdaBoost Validation Set Confusion Matrix

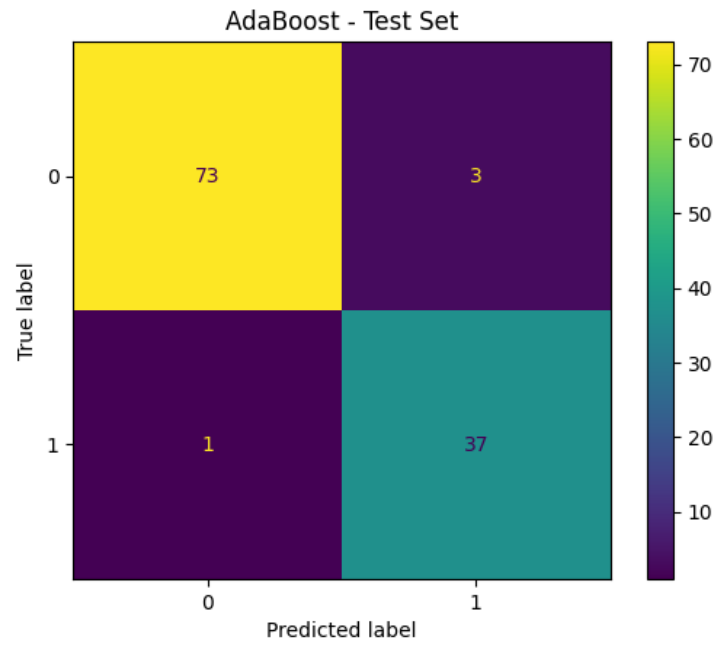


Figure 25: AdaBoost Test Set Confusion Matrix

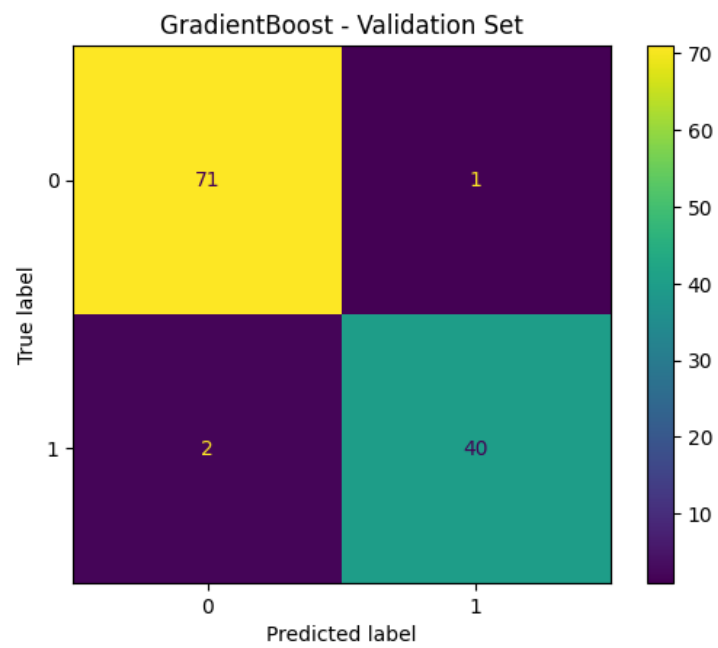


Figure 26: Gradient Boost Validation Set Confusion Matrix

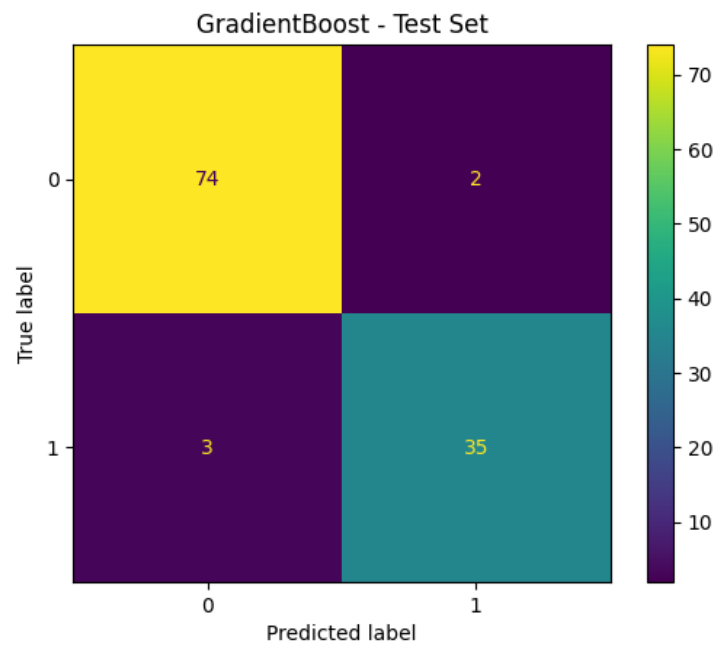


Figure 27: Gradient Boost Test Set Confusion Matrix

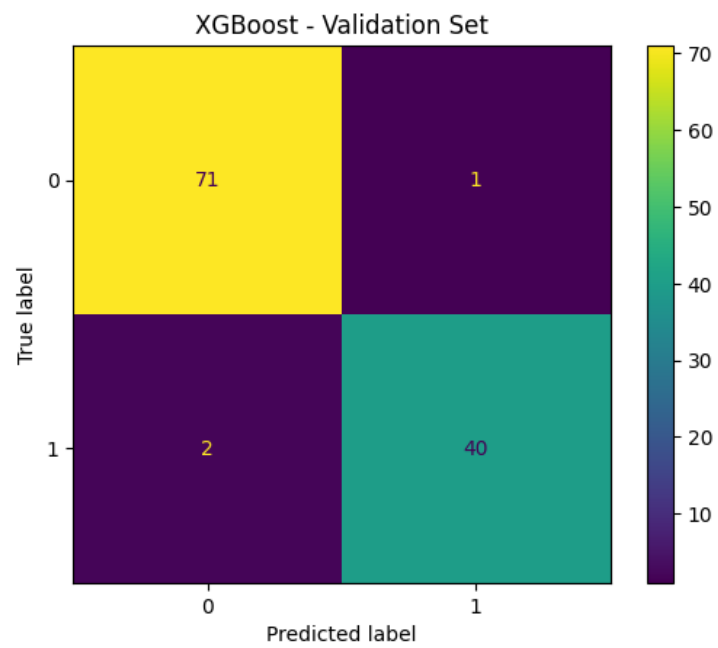


Figure 28: XGBoost Validation Set Confusion Matrix

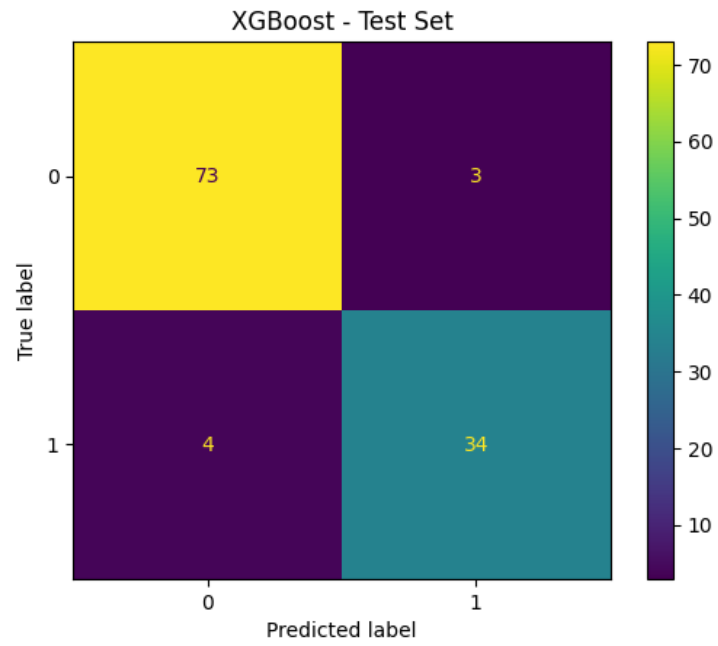


Figure 29: XGBoost Test Set Confusion Matrix

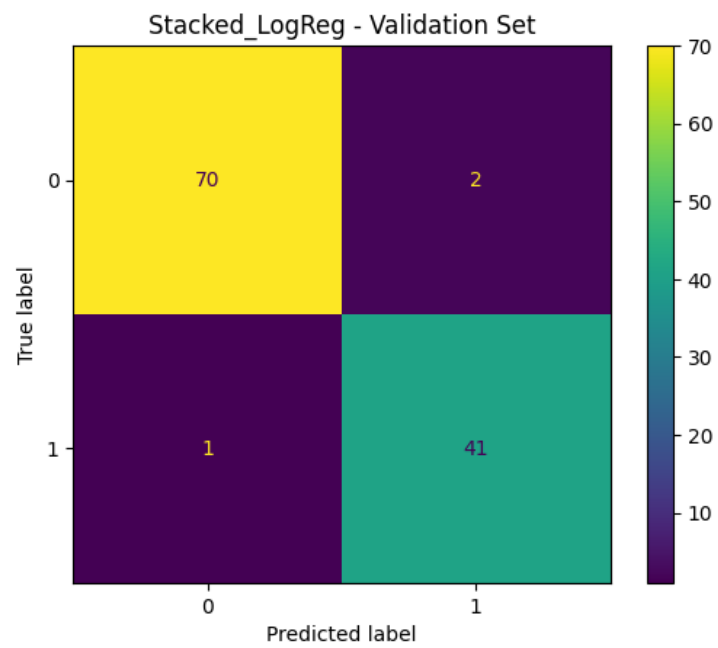


Figure 30: Stacked Ensemble 1 Validation Set Confusion Matrix

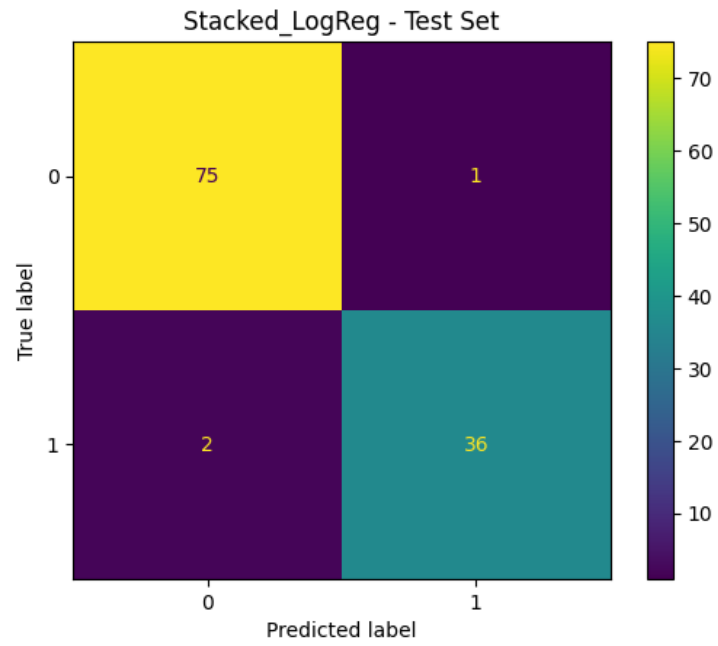


Figure 31: Stacked Ensemble 1 Test Set Confusion Matrix

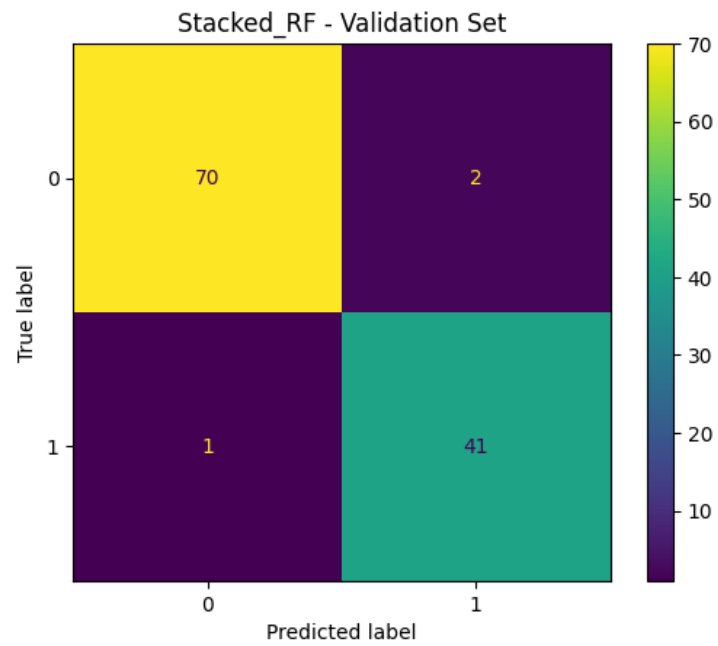


Figure 32: Stacked Ensemble 2 Validation Set Confusion Matrix

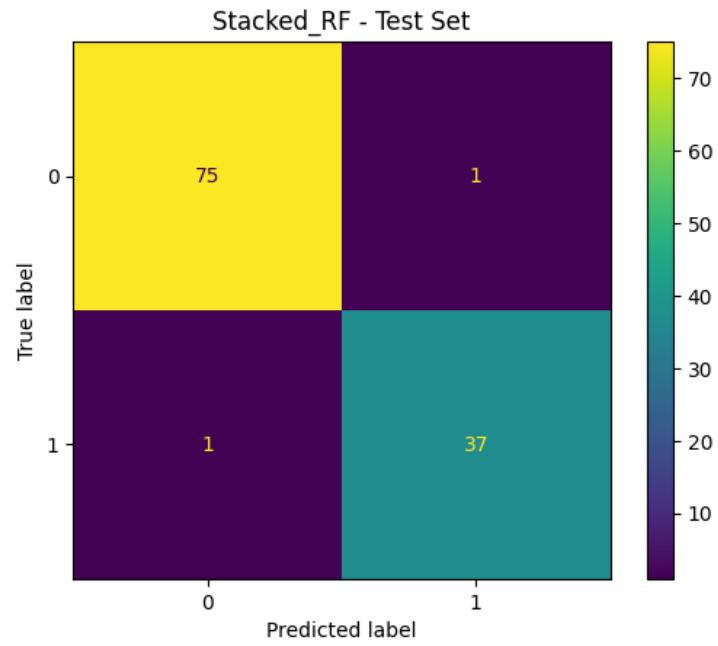


Figure 33: Stacked Ensemble 2 Test Set Confusion Matrix

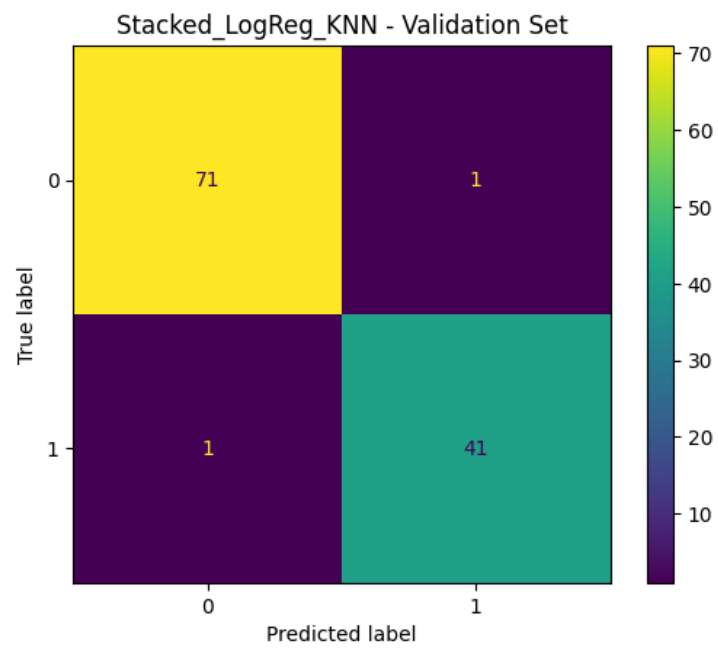


Figure 34: Stacked Ensemble 3 Validation Set Confusion Matrix

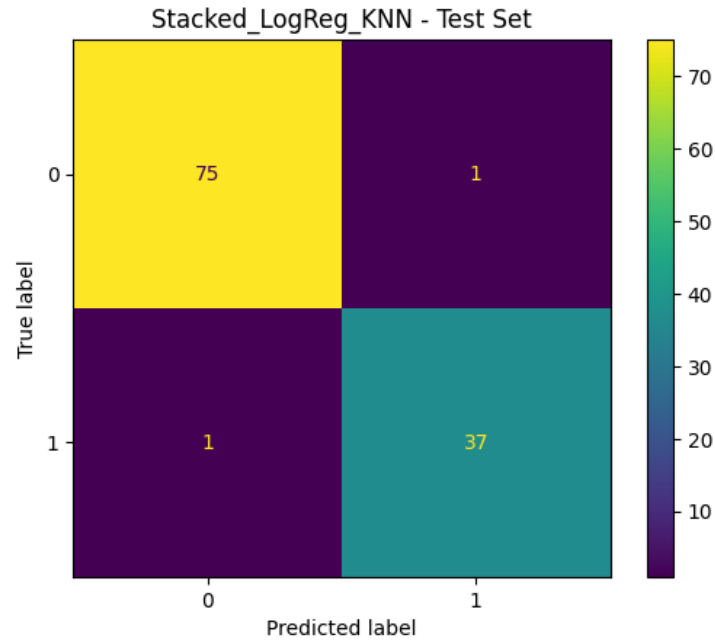


Figure 35: Stacked Ensemble 3 Test Set Confusion Matrix

Results

Table 1: Decision Tree - Hyperparameter Tuning

Criterion	Max Depth	Accuracy	F1 Score
gini	3	0.88	0.87
entropy	5	0.90	0.89
gini	None	0.91	0.90

Table 2: AdaBoost - Hyperparameter Tuning

n estimators	Learning rate	Accuracy	F1 Score
50	0.5	0.92	0.91
100	1.0	0.93	0.92
50	1.0	0.94	0.93

Table 3: Gradient Boosting - Hyperparameter Tuning

n estimators	Learning rate	Max Depth	Accuracy	F1 Score
100	0.05	3	0.93	0.92
200	0.1	2	0.94	0.93
100	0.1	3	0.95	0.94

Table 4: XGBoost - Hyperparameter Tuning

n estimators	Learning rate	Max Depth	Gamma	Accuracy	F1 Score
100	0.1	2	0.0	0.94	0.93
200	0.05	3	0.2	0.95	0.94
100	0.1	3	0.5	0.96	0.95

Table 5: Random Forest - Hyperparameter Tuning

n estimators	Max Depth	Criterion	Accuracy	F1 Score
100	10	gini	0.93	0.92
150	None	entropy	0.94	0.93
200	None	gini	0.95	0.94

Table 6: Stacked Ensemble - Hyperparameter Tuning

Base Models	Final Estimator	Accuracy / F1 Score
SVM, Naive Bayes, Decision Tree	Logistic Regression	0.97 / 0.97
SVM, Naive Bayes, Decision Tree	Random Forest	0.97 / 0.97
SVM, Decision Tree, KNN	Logistic Regression	0.98 / 0.98

Table 7: 5-Fold Cross Validation Results for All Models

Model	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Avg Accuracy
Decision Tree	0.90	0.91	0.92	0.90	0.91	0.91
AdaBoost	0.93	0.94	0.95	0.94	0.94	0.94
Gradient Boosting	0.94	0.95	0.96	0.95	0.94	0.95
XGBoost	0.95	0.96	0.96	0.95	0.96	0.96
Random Forest	0.94	0.95	0.95	0.96	0.95	0.95
Stacked Model	0.95	0.96	0.96	0.96	0.96	0.96

Observations and Conclusion

- Among all individual models, XGBoost achieved the best accuracy due to parallel tree construction and regularization (Ridge and Lasso penalties used).
- Ensemble methods outperformed the standalone Decision Tree as they prevent overfitting by using forward or backward propagation.
- Stacked models combining multiple classifiers along with a final estimator performed better than individual ensemble methods as they combine aggregate views of multiple models suitable for the dataset.

Best Practices and Learning Outcomes

- Learned to hyperparameter tune models using Grid Search.
- Verify the correctness of the best model chosen by printing their performances.
- Learned working of decision trees and ensemble techniques.
- Learned to create a stacked model with required individual models and a final estimator.
- Compared decision trees, bagging, boosting and stacked models for a classification dataset.